

AIOPS BOOTCAMP · GRADUATION PROJECT

EquiTeam

A fair, transparent, and explainable way to form our 5 project teams — balanced on gender, skill, and who works well with whom.

24

people

5

teams

10M / 14W

gender split

1 Why we're not just picking teams ourselves

We're 24 people about to spend weeks together building a real AIOps project — and right after this, we go into our OJT/internship placements. Whoever we're teamed with matters. Left to informal picking, teams naturally end up uneven: friend groups cluster together, the strongest people end up on one or two "super teams," and some people get left out of the process entirely. Nobody chooses that on purpose — it just happens when 24 people all try to pick at once.

The goal of EquiTeam is simple: every team should have a realistic, fair shot — not identical people, but comparable overall strength, a mix of skills, a fair gender split, and nobody stuck with someone who has openly said they can't work with them.

الهدف بسيط: كل فريق يأخذ فرصة عادلة — مش نفس الناس، لكن قوة متقاربة، ومهارات متنوعة، وتوزيع عادل للجنس، وحد محدش بيتخط مع حد رافض يشتغل معاه فعليًا.

2 How the ratings work — and why it's actually anonymous

Everyone fills out one short form rating *every other person* in the group, on two things:

1 Skill

Rated on 4 separate areas — DevOps, AI/ML, general technical ability, and project involvement — not one vague "how good are they" number.

قارة (٤ محاور)

2 Willingness to work together

Simply: would you want this person on your team? This is what protects against putting people who genuinely can't work together into the same team.

الرغبة في العمل سويا

Only the instructor sees individual responses. The algorithm only ever sees the computed averages per person — never "who said what about whom." Nobody, including whoever builds the tool, can trace a specific rating back to a specific rater in the final output.

الاستمارة فردية والدكتور بس اللي شايف الإجابات الأصلية، الخوارزمية شغالة على المتوسطات بس.

You also rate **yourself** on the same 4 skill areas. This isn't decorative — it's compared against how your peers rated you (Section 5 explains exactly why).

3 The five rules teams are built on — in order


These rules are checked in a strict order. An earlier rule is **never** broken to satisfy a later one.

#	Rule	What it means for you
1	Gender balance	Every team gets the same 2 men / mostly-3 women pattern. This is fixed before anything else is even calculated.
2	Skill spread	Every team gets a real mix — strong, mid, and developing skill levels — not "2 stars and 3 strugglers" hiding behind a matching average.
3	No forced pairing with someone who said no	If you specifically rated someone as "would not work with," the system routes around that pairing wherever possible — before teams are even finalized.
4	Minimum skill floor	No team's average skill can fall below an agreed line. This protects every team from being stuck without enough technical capability.
5	Minimum "getting along" floor	No team's <i>worst</i> internal relationship can be below an agreed comfort line — checked completely separately from skill.

Skill and willingness never cancel each other out. A team can't be shipped with bad skill just because everyone gets along great — and it can't be shipped with people who hate working together just because the skill numbers look good. Both have to pass, independently.

المهارة والرغبة في التعاون مش بيعوضوا عن بعض — لازم الاتنين يعدوا الحد الأدنى، كل واحد لوحده.

4 How a team actually gets built — step by step




Why "skill tiers" instead of just an average

Two teams can have the exact same *average* skill score while being completely different in practice — one team could be two very strong people and three people who need a lot of support, while another is five people of genuinely similar ability. Same average, very different reality. So instead of just matching totals, every skill pool (men and women, separately) is split into tiers from strongest to weakest, and **every team gets exactly one person from every tier** — guaranteeing a real spread, not a lucky-looking average.

WORKED EXAMPLE — 5 TEAMS, 10 MEN

Men are ranked 1 (strongest) through 10 (needs the most support) and split into 5 tiers of 2. Each team draws exactly one man from each tier — so every team ends up with one of the two strongest men, one of the two weakest, and three from the tiers in between. The same happens independently for the 14 women. No team can end up with "the two strongest men" together, because each tier only ever sends one person to each team.



5 What happens if you genuinely can't work with someone

An overall "getting along" score for a team can hide one bad relationship behind three good ones. So there's a separate, stronger rule just for this: if you rate someone specifically low on willingness — a real "I would actively avoid working with this person," not just a lukewarm rating — the system treats that as a **veto** and tries to seat you and that person on different teams entirely, before teams are even finalized.

This is checked at the moment people are being placed onto teams, not fixed afterward. It looks for another same-gender person with a similar skill level to place instead — so avoiding a bad pairing doesn't come at the cost of breaking the gender or skill balance.

الرفض المحدد بين شخصين معينين ينفادى وهو يبينى الفرق، مش بعد ما تتبنى وتتصلح.

In rare cases with a small group, it may not be possible to avoid every single veto without breaking gender or skill balance. If that happens, it's flagged explicitly for the instructor to make the final call — the algorithm never silently ignores it.

6 Making the skill score fair to quiet people — and hard to game

Peer ratings alone have a real weakness: confident, visible people tend to get rated higher regardless of actual ability, while quieter or more modest people get underrated. Adding a self-rating helps fix that — but it opens an obvious new problem: what stops someone from rating themselves artificially low to land on an "easier" team?

The fix: compare what you say about yourself to what your peers say about you

Your final skill score for each of the 4 areas is mostly based on your peers' average rating of you (80%), with your own self-rating included at a smaller weight (20%) — enough to help correct for being underrated, not enough to let you set your own placement.

WHAT HAPPENS IF THE NUMBERS DON'T MATCH

If your self-rating and your peer average are far apart in either direction — you rated yourself much lower *or* much higher than your peers did — your self-rating is dropped entirely for that specific skill area, and your peer average is used instead. Both directions are checked, so there's no "safe direction" to lie in: trying to look weaker to land on an easier team, or stronger to land on a tougher one, has the same result — your own number gets ignored and the group's honest assessment is used.

This is about catching numbers that don't add up — not about accusing anyone of lying. Someone might genuinely be modest, or genuinely underrated by peers who don't know their work well. Either way, the fix is the same: lean on the group's collective view, and flag the mismatch for the instructor to see.

7 What you'll actually see at the end

Once teams are formed, a report goes to the instructor (not the raw scores) showing:

- The final 5 teams, with each team's overall skill balance and gender split.
- Any place where the algorithm had to make a trade-off, and exactly why.
- Anything it couldn't fully resolve — flagged clearly for a human decision, never hidden or guessed at.

Nothing here is a black box. Every team can be explained: "you're on this team because of your skill tier, your gender slot, and because no one on it vetoed working with you." If anyone asks why the teams turned out the way they did, there's a real, specific answer — not "the algorithm decided."

مفیش قرار غامض — أي فريق ممكن نوضح بالطبط ليه اتكون كده.

8 Questions people usually ask

Q: Can I end up on a team with someone I dislike?

Only if you never actually flagged them as a hard "would not work with" — ordinary lukewarm feelings aren't enough to trigger the veto rule, but a real, specific rejection is respected and actively avoided.

Q: What if I'm just not confident rating people's skill?

That's exactly why it's split into 4 specific areas instead of one vague number, and why your rating is only part of the picture — it's averaged with everyone else's view of that person, so no single rating decides anyone's fate.

Q: Could someone game this by rating their friends high and everyone else low?

The system checks for raters whose skill and willingness scores move suspiciously together for everyone they rated — a sign of rating based on friendship rather than honest assessment. This gets flagged for the instructor to see, not silently ignored.

Q: Will the strongest people all end up on one team?

No — this is exactly what the tier system in Section 4 prevents. Every team gets one person from every skill tier, so it's mathematically impossible for two people from the top tier to end up on the same team.

Q: Who makes the final call if something can't be resolved?

The instructor. The algorithm never guesses through a genuine conflict — anything it can't cleanly resolve is reported explicitly, with the specific reason, for a human decision.